

# Model evaluation and cross-validation

German Demidov

Institut für Medizinische Genetik und Angewandte Genomik, Tübingen

*german.demidov at med.uni-tuebingen.de*

September 16, 2022

# Overview

- 1 Introduction
- 2 Model evaluation
- 3 Cross-validation

# Twitter Mixup



Solve-RD  
@Solve\_RD

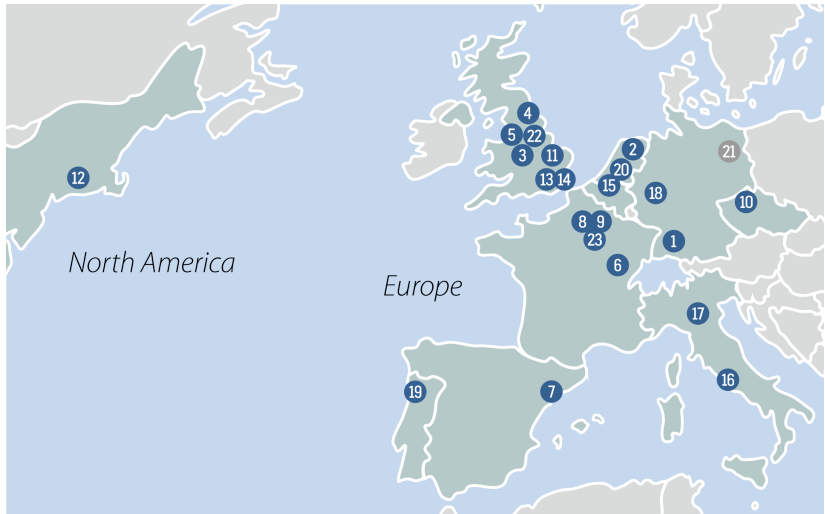


Tomorrow, our partner German Demidov [@not\\_a\\_reptiloid](#) from [@uni\\_tue](#) [@uktuebingen](#) [@MedTuebingen](#) will present Solve-RD & speak about “Cross-validation & model evaluation” at the [#NGSchool22](#) in Jablonna 🇵🇱 Full programme [👉 ngschool.eu/ngschool2022/](#)

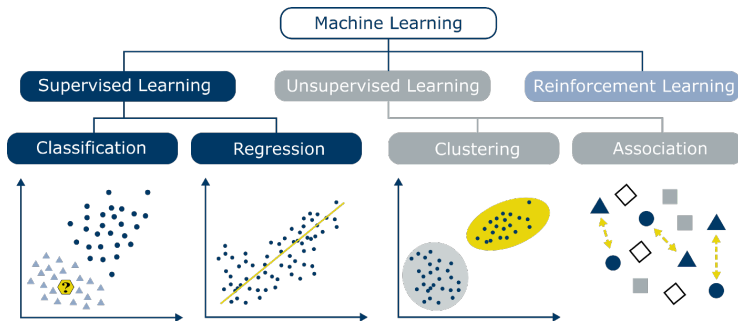
[@NGSchoolEU](#) [#bioinformatics](#)  
[#NGS](#)

12:19 PM · Sep 16, 2022 · Twitter Web App

# SolveRD



# Machine learning problems



# Machine learning model

Model:

$$y \sim M(X)$$

Real data:  $\{y_i, x_{i1}, \dots, x_{ip}\}_{i=1}^n$

# Evaluation metric

A way to quantify the performance of a machine learning model

We need it to 1) compare with other models, 2) select the best hyperparameters for our model

(Disclaimer: everything which is in this presentation I copied from somewhere: Introduction to Stat Learning, youtube lectures of M. Khalusova, MarinStatsLectures and others, personal blog of Fabio Sigrüst and others I found on Google)

# Linear regression example

$$\mathbf{y} = X\boldsymbol{\beta} + \boldsymbol{\epsilon},$$

where

$$\mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix}, \quad \mathbf{X} = \begin{pmatrix} \mathbf{x}_1^T \\ \mathbf{x}_2^T \\ \vdots \\ \mathbf{x}_n^T \end{pmatrix} = \begin{pmatrix} 1 & x_{11} & \cdots & x_{1p} \\ 1 & x_{21} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & \cdots & x_{np} \end{pmatrix},$$

$$\boldsymbol{\beta} = \begin{pmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{pmatrix}, \quad \boldsymbol{\epsilon} = \begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{pmatrix}.$$



# Types of problems and Metrics

$$y \sim M(X)$$


- Regression (continuous outcome):  $R^2$ , MAE, RMSE
- Ordinal outcome (RMSE)
- Classification (factor outcome): accuracy, Precision, Recall, F1 Score, ROC/AUC, Precision/Recall AUC, Matthews correlation coefficient

# Machine learning vs statistics

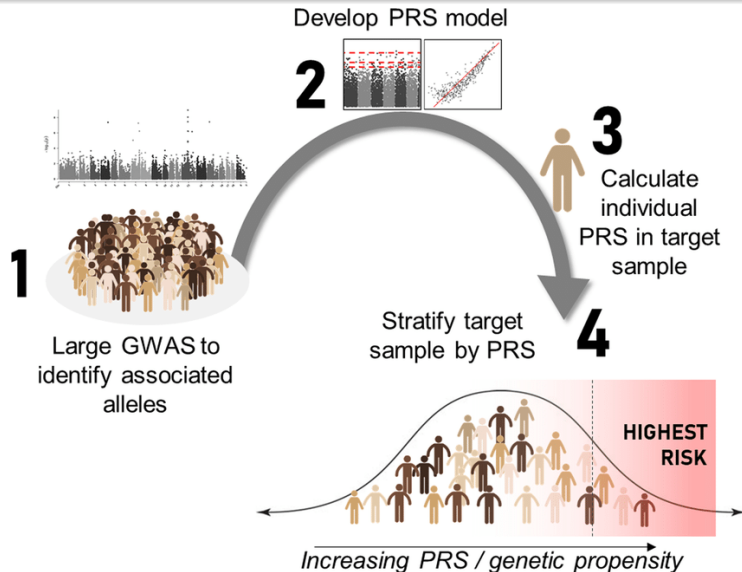
## Machine Learning




## Statistics



# Inference vs Prediction



## Difference between inference vs Prediction: Confounding




Q: When it is important?

- Estimation of the effect size
- Prediction

# Types of error


- Train error
- Test error



# Evaluation of a regression model

- Coefficient of Determination ( $R^2$ )
- RMSE (root mean squared error)
- MAE (mean absolute error)

# Residuals



# Coefficient of Determination

We calculate the mean of our outcome:

$$\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$$

The sum of squares of residuals, also called the residual sum of squares:

$$SS_{\text{res}} = \sum_i (y_i - f_i)^2 = \sum_i e_i^2$$

The total sum of squares (proportional to the variance of the data):


$$SS_{\text{tot}} = \sum_i (y_i - \bar{y})^2$$

The most general definition of the coefficient of determination is

$$R^2 = 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}$$







# Coefficient of Determination: Graphical Explanation



$$R^2 = 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}$$

In linear least squares multiple regression with an intercept term,  $R^2$  = the square of the Pearson correlation coefficient between the observed  $y$  and modeled (predicted)  $f$  data values of the dependent variable.

# Drawbacks



# Multiple linear regression

```
set.seed(1)
outcome <- rnorm(100)
predictor <- rnorm(100)
newlm <- lm(outcome ~ predictor)
summary(newlm)
```

# Multiple linear regression

Call:

```
lm(formula = outcome ~ predictor)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.32416	-0.60361	0.00536	0.58305	2.29316

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.1088521	0.0903480	1.205	0.231
predictor	-0.0009324	0.0947216	-0.010	0.992

Residual standard error: 0.9028 on 98 degrees of freedom

Multiple R-squared: 9.887e-07, Adjusted R-squared: -0.0102


F-statistic: 9.689e-05 on 1 and 98 DF, p-value: 0.9922

# Multiple linear regression

```
R_squared <- c(summary(newlm)$r.squared)
for (i in 1:50) {
  new_predictor <- rnorm(100)
  predictor = cbind(predictor, new_predictor)
  newlm <- lm(outcome ~ predictor)
  R_squared <- c(R_squared, summary(newlm)$r.squared)
}
plot(R_squared)
```

## $R^2$ with more random predictors

$R^2$  alone cannot be used as a meaningful comparison of models with very different numbers of independent variables. F-test can be performed on the residual sum of squares.



# Adjusted $R^2$

Replace r.squared with adj.r.squared:

```
R_squared <- c(R_squared, summary(newlm)$adj.r.squared)
```

$$\bar{R}^2 = 1 - (1 - R^2) \frac{n - 1}{n - p}$$

where  $p$  is the total number of explanatory variables in the model, and  $n$  is the sample size.

# Model selection via BIC and AIC

Not model evaluation, but selection

$$\text{AIC} = 2k - 2\ln(\hat{L})$$

$$\text{BIC} = k \ln(n) - 2\ln(\hat{L})$$

where  $k$  is the number of estimated parameters,  $n$  is the number of observations.



## Explained variance in Random Forest regression

```
library(randomForest)
library(ggplot2)
dat <- data.frame(ggplot2::diamonds[1:1000,1:7])
rf <- randomForest(formula = carat ~ .,
data = dat, ntree = 500)
rf
# Call:
#   randomForest(formula = carat ~ ., data = dat, ntree = 500)
#
#           Type of random forest: regression
#           Number of trees: 500
# No. of variables tried at each split: 2
# Mean of squared residuals: 0.001820046
# % Var explained: 95.22
```


## Manual $R^2$ for random forest regression models

```
actual = dat$carat  
predicted <- unname(predict(rf, dat))  
  
R2 <- 1 - (sum((actual-predicted)^2)/sum(  
  (actual-mean(actual))^2))
```

## $R^2$ using test data

```
library(faux)
num_of_predictors <- 100
bvn <- rnorm_multi(200, num_of_predictors, 0, 1, .75)
colnames(bvn)[1] = "a"
r2 <- c()
for (i in 2:num_of_predictors) {
  train <- bvn[1:100,1:i]
  test <- bvn[101:200,1:i]
  newlm <- lm(a ~ ., data=train)
  predictions <- predict(newlm, test)
  r2 <- c(r2, cor(predictions, test$a)**2)
}
plot(r2)
```


# $R^2$ using test data



## $R^2$ using test data

```
library(randomForest)
r2 <- c()
for (i in 2:num_of_predictors) {
  train <- bvn[1:100,1:i]
  test <- bvn[101:200,1:i]
  randf <- randomForest(a ~ ., data=train)
  predictions <- predict(randf, test)
  r2 <- c(r2, cor(predictions, test$a)**2)
}
plot(r2)
```

## $R^2$ using test data



# The use of coefficient of determination

R squared has an intuitive scale and does not depend of y unit  
R squared gives you no information about the prediction error

# RMSE and MAE

$$RMSE(y, \hat{y}) = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

- most commonly used metric


Should Mean Absolute Error be used instead?

$$MAE = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{n}$$




# RMSE and MAE and other losses


a) Squared errors




b) Absolute errors



c) Winsorized



d) Bisquare



# MAE and RMSE



In common:

- Same units as the  $y$  response values
- Don't care about the direction of errors

Different:

- Higher weights to large errors in RMSE
- RMSE is differentiable

# RMSE application



# Bias and variance


$D$  is the training dataset

$$\mathbb{E}_{D,\epsilon} \left[ (y - \hat{f}(x; D))^2 \right] = \left( \text{Bias}_D [\hat{f}(x; D)] \right)^2 + \text{Var}_D [\hat{f}(x; D)] + \sigma^2$$


Variance refers to the amount by which  $\hat{f}$  would change if we estimated it using a different training data set.

Bias refers to the error that is introduced by approximating a real-life problem, which may be extremely complicated, by a much simpler model.

# Bias and variance








# Bias and variance



**FIGURE 2.10.** Details are as in Figure 2.9, using a different true  $f$  that is much closer to linear. In this setting, linear regression provides a very good fit to the data.

# Classification

	Feature 1	Feature 2	Feature 3	Feature 4	...	Feature n	Classes
Instance 1	Land	Legs	Mammal	Barks	...	n1	 dog
Instance 2	Land	Legs	Mammal	Purrs	...	n2	 cat
Instance 3	Ocean	Flippers	Mammal	Sings	...	n3	 whale
Instance 4	Ocean	Flippers	Fish	Silent	...	n4	 shark
Instance 5	Ocean	Claws	Crustacean	Silent	...	n5	 lobster

# Classification

$$\text{Accuracy} = \frac{\text{Number of correct prediction}}{\text{Total number of predictions}}$$

Is accuracy of 99% good?



## Confusion matrix

```
library(caret)
library(InformationValue)
library(ISLR)

data <- Default # Credit Card Default Data
set.seed(1)
sample <- sample(c(TRUE, FALSE), nrow(data),
  replace=TRUE, prob=c(0.7,0.3))
train <- data[sample, ]
test <- data[!sample, ]

model <- glm(default~student+balance+income,
  family="binomial", data=train)
```

# Confusion matrix

```
predicted <- predict(model, test, type="response")
#convert defaults from "Yes" and "No" to 1's and 0's
test$default <- ifelse(test$default=="Yes", 1, 0)
#find optimal cutoff probability to use to maximize accuracy
optimal <- optimalCutoff(test$default, predicted)[1]
#create confusion matrix
confusionMatrix(test$default, predicted //def thresh 0.5
  0  1
0 2912 64
1  21 39
confusionMatrix(test$default, predicted, threshold=optimal)
  0  1
0 2919 68
1  14 35
```

# Confusion matrix

```
optimal <- optimalCutoff(test$default, predicted,  
optimiseFor="Ones")[1]
```

```
confusionMatrix(test$default, predicted,  
threshold=optimal)
```

	0	1
0	2003	3
1	930	100

# Precision, Recall

		Predicted condition		Sources: [6][7][8][9][10][11][12][13][14] view-talk-edit	
Actual condition	Total population $= P + N$	Positive (PP)	Negative (PN)	Informedness, bookmaker informedness (BM) $= TPR + TNR - 1$	Prevalence threshold (PT) $= \frac{\sqrt{TPR \times FPR} - FPR}{TPR - FPR}$
	Positive (P)	True positive (TP), hit	False negative (FN), type II error, miss, underestimation	True positive rate (TPR), recall, sensitivity (SEN), probability of detection, hit rate, power $= \frac{TP}{P} = 1 - FNR$	False negative rate (FNR), miss rate $= \frac{FN}{P} = 1 - TPR$
	Negative (N)	False positive (FP), type I error, false alarm, overestimation	True negative (TN), correct rejection	False positive rate (FPR), probability of false alarm, fall-out $= \frac{FP}{N} = 1 - TNR$	True negative rate (TNR), specificity (SPC), selectivity $= \frac{TN}{N} = 1 - FPR$
	Prevalence $= \frac{P}{P + N}$	Positive predictive value (PPV), precision $= \frac{TP}{PP} = 1 - FDR$	False omission rate (FOR) $= \frac{FN}{PN} = 1 - NPV$	Positive likelihood ratio (LR+) $= \frac{TPR}{FPR}$	Negative likelihood ratio (LR-) $= \frac{FNR}{TNR}$
	Accuracy (ACC) $= \frac{TP + TN}{P + N}$	False discovery rate (FDR) $= \frac{FP}{PP} = 1 - PPV$	Negative predictive value (NPV) $= \frac{TN}{PN}$ $= 1 - FOR$	Markedness (MK), deltaP ( $\Delta p$ ) $= PPV + NPV - 1$	Diagnostic odds ratio (DOR) $= \frac{LR+}{LR-}$
	Balanced accuracy (BA) $= \frac{TPR + TNR}{2}$	F <sub>1</sub> score $= \frac{2PPV \times TPR}{PPV + TPR} = \frac{2TP}{2TP + FP + FN}$	Fowlkes–Mallows index (FM) $= \sqrt{PPV \times TPR}$	Matthews correlation coefficient (MCC) $= \sqrt{TPR \times TNR \times PPV \times NPV}$ $= \sqrt{TPR \times TNR \times PPV \times NPV}$ $= \sqrt{TPR \times TNR \times PPV \times NPV}$	Threat score (TS), critical success index (CSI), Jaccard index $= \frac{TP}{TP + FN + FP}$

# Precision, Recall

Precision =  $\frac{TP}{TP+FP}$  - concentrate on True Positives

Recall =  $\frac{TP}{TP+FN}$  - concentrate on False Negatives

What's more important to minimize? You decide

# F1

Harmonic mean between precision and recall

$$F1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} = \frac{2 \cdot TP}{2 \cdot TP + FP + FN}$$

# Matthews Correlation Coefficient

$$MCC = \frac{TP \cdot TN - FP \cdot FN}{\sqrt{(TP+FP)(TP+FN)(TN+FP)(TN+FN)}}$$

Takes into account all four components

# MCC vs F1

	0	1
0	TN=0	FP=5
1	FN=0	TP=95

Accuracy: 95%

F1 score = 0.974

MCC = Undefined (dividing by 0)



# MCC vs F1

	0	1
0	TN=1	FP=4
1	FN=5	TP=90

Accuracy: 95%

F1 score = 0.952

MCC = 0.135

# MCC vs F1

	0	1
0	TN=90	FP=5
1	FN=4	TP=1

Accuracy: 95%


F1 score = 0.182

MCC = 0.135

# MCC vs F1

MCC can be better to summarize the confusion matrix (only for binary problems)

# ROC / AUC



# How to calculate ROC / AUC in R


Packages: pROC, ROCR

<https://www.r-bloggers.com/2016/11/calculating-auc-the-area-under-a-roc-curve/>

# How to calculate ROC / AUC in R

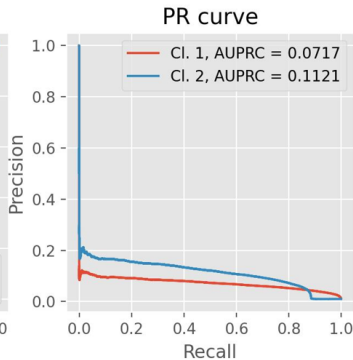

```
library(ROCR)
data(ROCR.simple)
df <- data.frame(ROCR.simple)
pred <- prediction(df$predictions, df$labels)
perf <- performance(pred,"tpr","fpr")
library(pROC)
pROC_obj <- roc(df$labels,df$predictions, smoothed = TRUE,
               # arguments for ci
               ci=TRUE, ci.alpha=0.95, stratified=FALSE,
               # arguments for plot
               plot=TRUE, auc.polygon=TRUE, max.auc.polygon=TRUE,
               print.auc=TRUE, show.thres=TRUE)
sens.ci <- ci.se(pROC_obj)
plot(sens.ci, type="shape", col="lightblue")
plot(sens.ci, type="bars")
```

# How to calculate ROC / AUC in R



# Precision / Recall curve


2 classifiers and imbalanced data





# AUC ROC in imbalanced data


ROC curve



(Sin-Yi Chou github)

# AUC PR in imbalanced data

PR curve



# AUC PR vs AUC ROC

- ROC in case the model is biased towards one class can't be large - unlike Accuracy. If the goal of the model is to perform equally well on both classes, ROC is the choice
- In case you care about the small positive class, PR AUC can be better
- If the model doesn't work after the metric is changed, there are still other remedies to deal with imbalanced data, such as downsampling/upsampling

# Multiple class: Precision

	Bird	Cat	Dog
Bird	1	0	1
Cat	0	4	0
Dog	0	1	2

	TP	FP	Pr	N samples
Bird	1	0	1	2
Cat	4	1	0.8	4
Dog	2	1	0.66	3
Total	7	2		

$$\text{Microprecision} = \frac{7}{7+2} = 0.7777$$


$$\text{Macroprecision} = \frac{1}{3}(1 + 0.8 + 0.66) = 0.822$$

$$\text{Weighted precision} = \frac{1 \cdot 2 + 0.8 \cdot 4 + 0.66 \cdot 3}{2 + 4 + 3} = 0.8$$

## Which to choose

- Microprecision: every sample contributed
- Macroprecision: all the classes equally contributed (good for underrepresented classes)
- Weighted: each classes' contribute to the precision according to the size of the class

# Train-test approach



**Model evaluation:** hold-out a test set

- `mod.fit(X_train, y_train)`
- `mod.predic(X_test)`




**Model selection:** hold-out a validation set

- Explore hyper-parameters grid
- Fit on train, evaluate on validation
- Pick best hyper-parameter



**Cross-validation:** when dataset is too small for to apply hold-out strategy. CV Can be used for model evaluation or/and model selection.



# Resampling methods

- Estimation: bootstrap, permutation-based tests
- Model evaluation: cross-validation

# Cross-validation types

- Leave one out CV
- k-fold CV
- repeated k-fold CV



# Summary of the error

$$CV_k = \frac{1}{k} \sum_{i=1}^k MSE_i$$


## Cross-validation types

Which cross-validation is better?

- LOOCV can take incredibly long, 10-fold can be faster
- LOOCV has almost the lowest bias (almost all the samples used for the training)
- LOOCV has a very high variance

Typically one performs k-fold cross-validation using  $k = 5$  or  $k = 10$  since empirically it shows good balance between variance and bias

# Cross-validation as parts of methods



*lambda.min* - value of lambda that gives minimum cvm.

*lambda.1se* - largest value of lambda such that error is within 1 standard error of the minimum.

# Cross-validation in caret (R)

```
set.seed(1)
data("swiss")
library(caret)
```

## Cross-validation in caret (LOOCV)

```
set.seed(1)
train.control <- trainControl(method = "LOOCV")
# Train the model
model <- train(Fertility ~., data = swiss, method = "lm",
               trControl = train.control)
print(model)
Linear Regression
47 samples
  5 predictor
No pre-processing
Resampling: Leave-One-Out Cross-Validation
Summary of sample sizes: 46, 46, 46, 46, 46, 46, ...
Resampling results:
```

RMSE	Rsquared	MAE
7.738618	0.6128307	6.116021

## Cross-validation in caret (K-fold)

```
set.seed(1)
train.control <- trainControl(method = "cv", number = 10)
# Train the model
model <- train(Fertility ~., data = swiss, method = "lm",
               trControl = train.control)
print(model)
Linear Regression
47 samples
  5 predictor
No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 42, 42, 43, 41, 42, 44, ...
Resampling results:
```

RMSE	Rsquared	MAE
7.464863	0.7227897	5.953467

# The End